

# ackstorm

**Kubernetes Training**  
2019 / Barcelona





**kubernetes**

# Advanced Controllers



- 1.- ReplicationController: A ReplicationController ensures that a specified number of pod replicas are running at any time. ReplicationController is one of the first controllers. It now is deprecated, and has been replaced by ReplicaSet.
- 2.- ReplicaSet is the next-generation Replication Controller. The only difference between a ReplicaSet and a Replication Controller is the selector support. ReplicaSet supports set-based selector requirements, whereas a Replication Controller only supports equality-based selector requirements.
- 3.- Deployments are the default way of creating pods. Using any other type of controller will require use case specifications. Just like a ReplicaSet, it ensures that a specified number of pod replicas are running at any time.

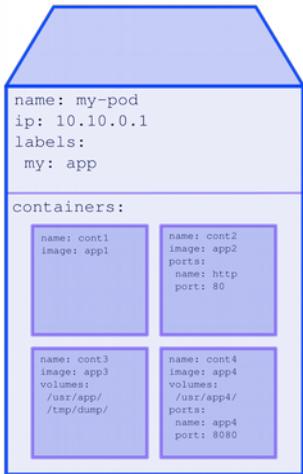
Note: Deployments do not create pods. They rather create ReplicaSets that create the pods.

- 4.- StatefulSets are objects used to manage stateful applications. They manage the deployment and scaling of a set of Pods, and provides guarantees about the ordering and uniqueness of these Pods; that are created from the same spec, but are not interchangeable since each has a persistent identifier that it maintains across any rescheduling.
- 5.- A DaemonSet ensures that all (or some) Nodes run a copy of a Pod. As nodes are added to the cluster, Pods are added to them. As nodes are removed from the cluster, those Pods are garbage collected.
- 6.- A job creates one or more pods and ensures that a specified number of them successfully terminate. As pods successfully complete, the job tracks the successful completions. When a specified number of successful completions is reached, the job itself is complete.
- 7.- A CronJob creates Jobs on a time-based schedule. A CronJob object is like one line of a crontab (cron table) file. It runs a job periodically on a given schedule, written in Cron format. A CronJob creates a job object about once per execution time of its schedule.

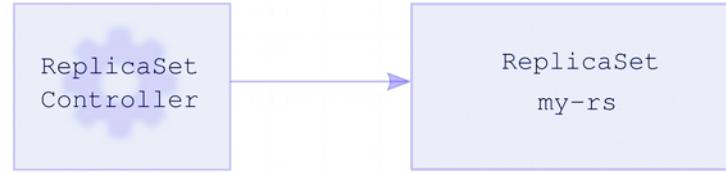
# ReplicaSet



```
name: my-rs  
replicas: 3  
Template:
```



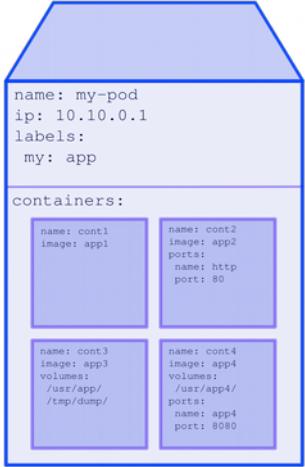
Cluster



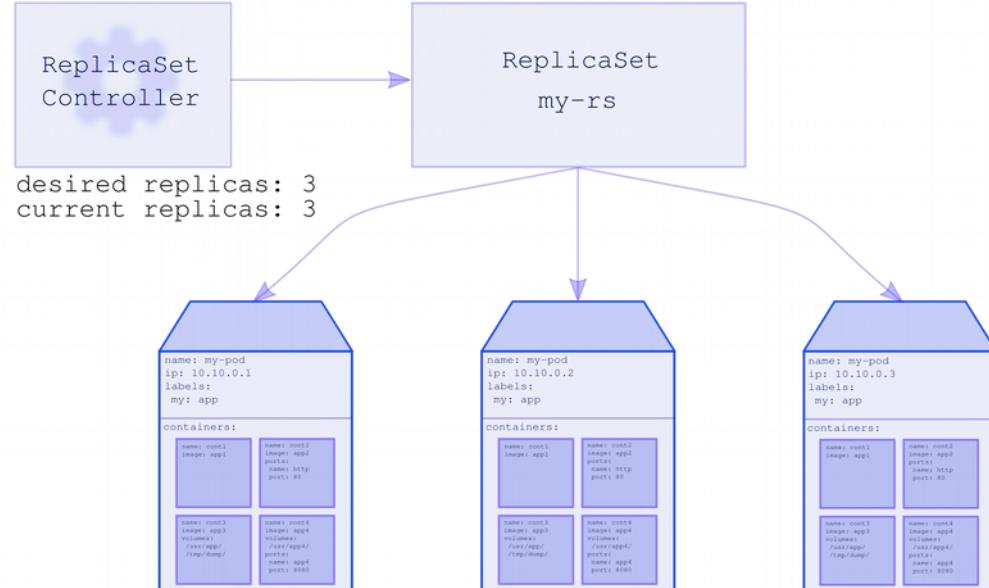
# ReplicaSet



```
name: my-rs  
replicas: 3  
Template:
```



## Cluster



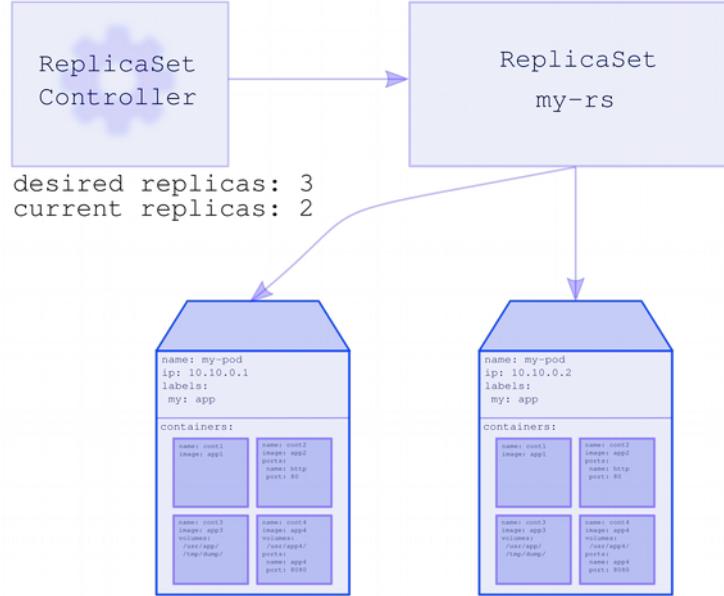
# ReplicaSet



```
name: my-rs  
replicas: 3  
Template:
```



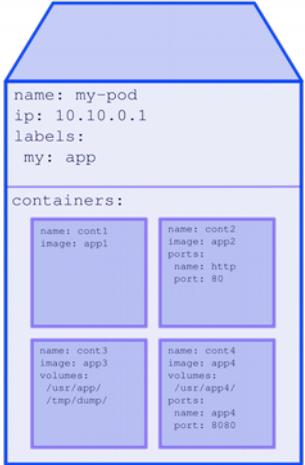
## Cluster



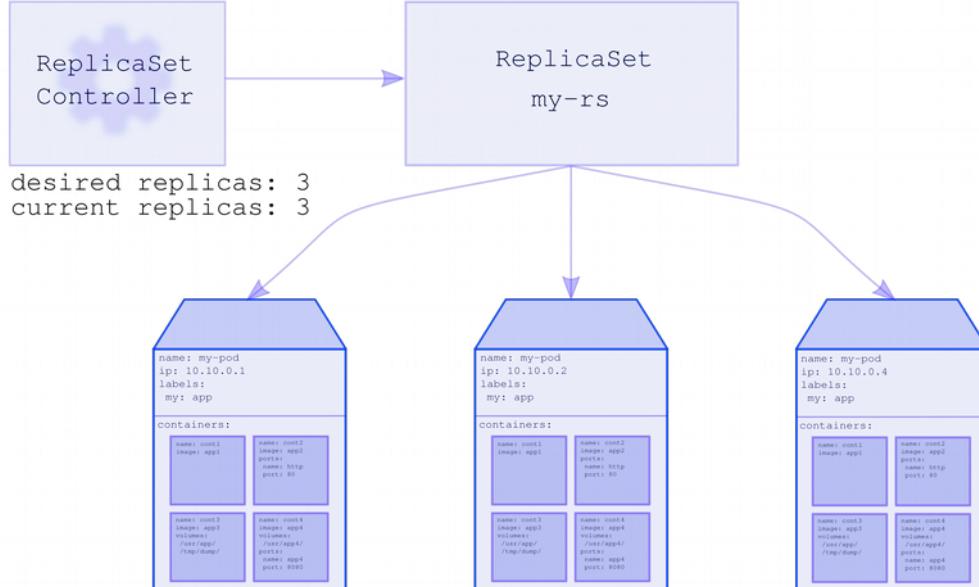
# ReplicaSet



```
name: my-rs  
replicas: 3  
Template:
```



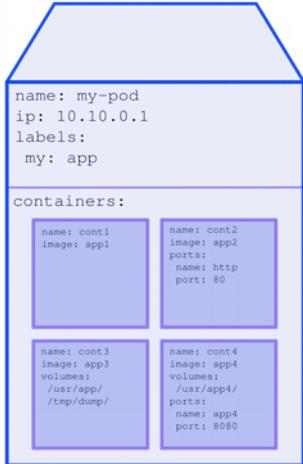
## Cluster



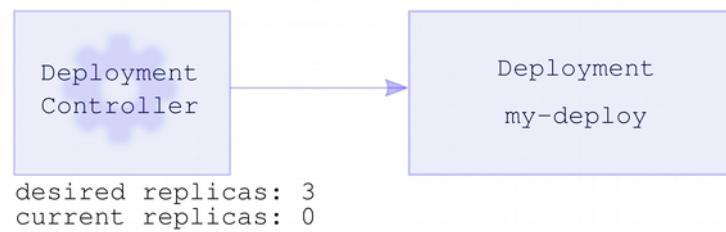
# Deployment



```
name: my-deploy  
replicas: 3  
strategy: RollingUpdate  
Template:
```



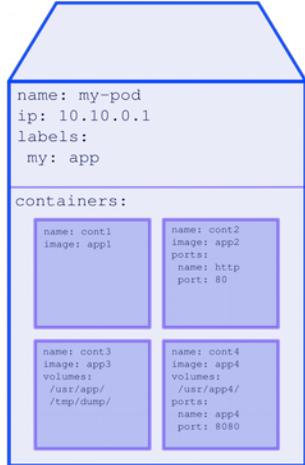
Cluster



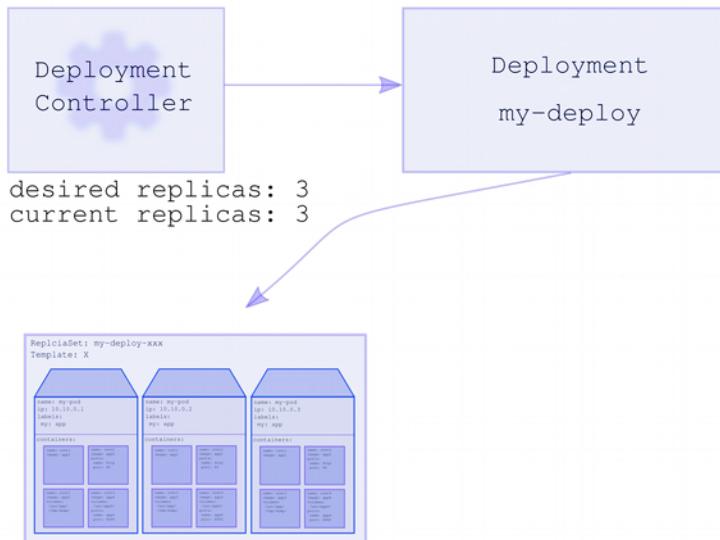
# Deployment



```
name: my-deploy  
replicas: 3  
strategy: RollingUpdate  
Template:
```



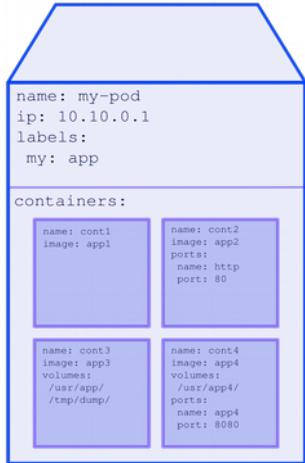
## Cluster



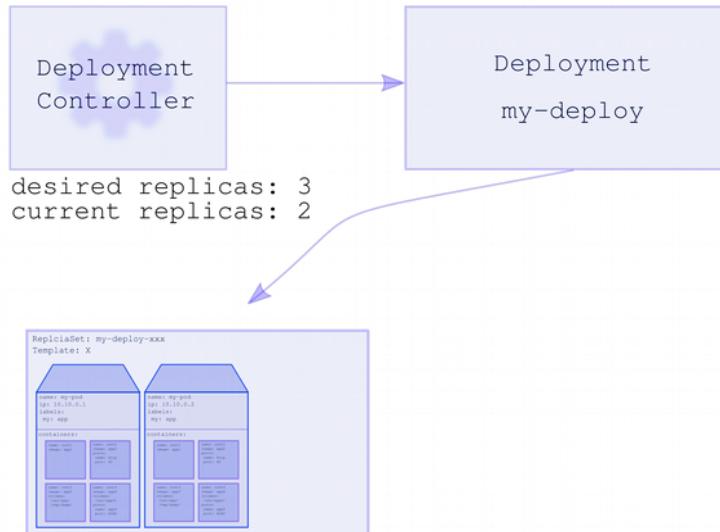
# Deployment



```
name: my-deploy  
replicas: 3  
strategy: RollingUpdate  
Template:
```



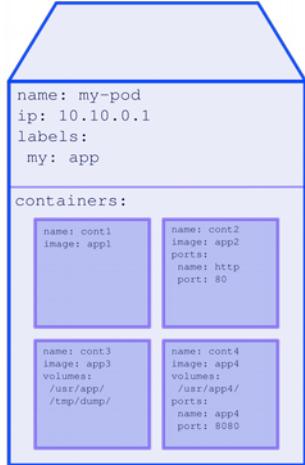
## Cluster



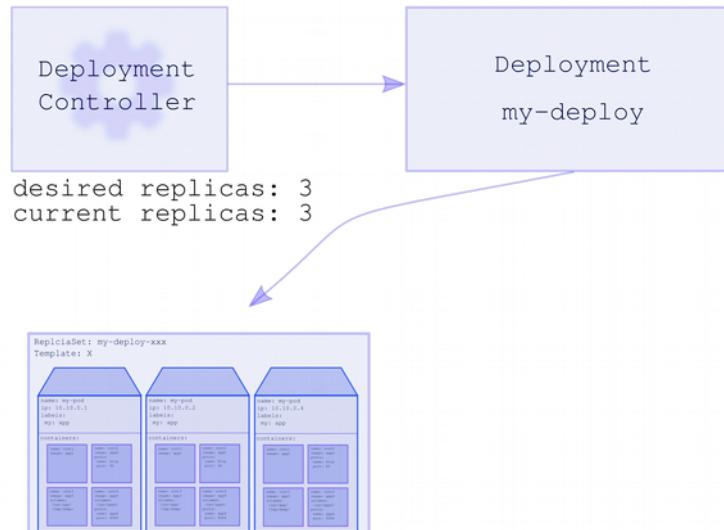
# Deployment



```
name: my-deploy  
replicas: 3  
strategy: RollingUpdate  
Template:
```



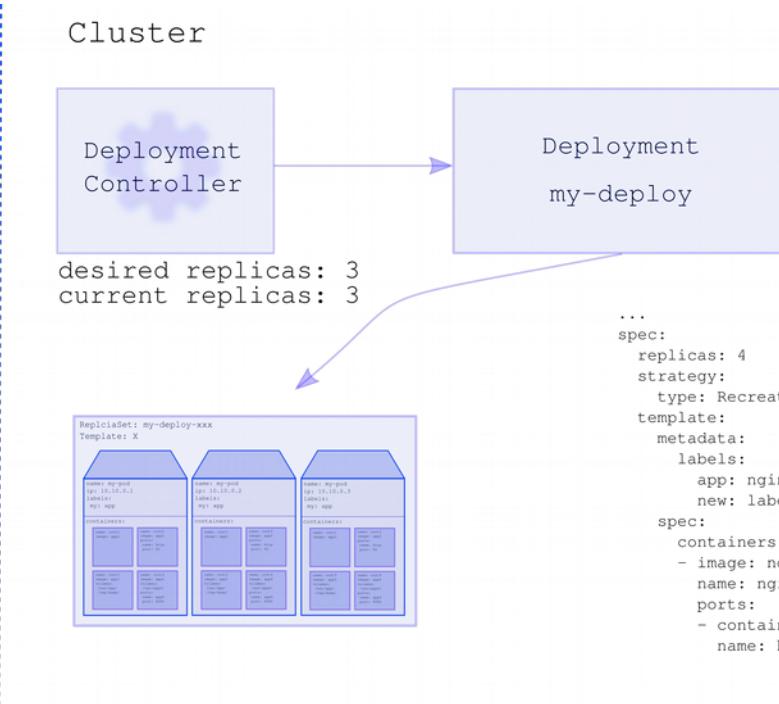
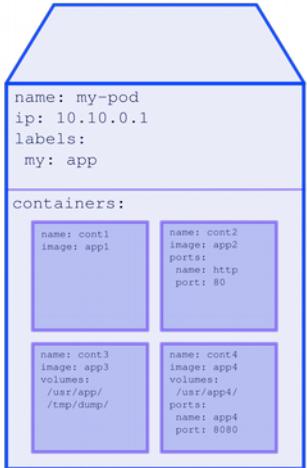
## Cluster



# Deployment



```
name: my-deploy  
replicas: 3  
strategy: RollingUpdate  
Template:
```

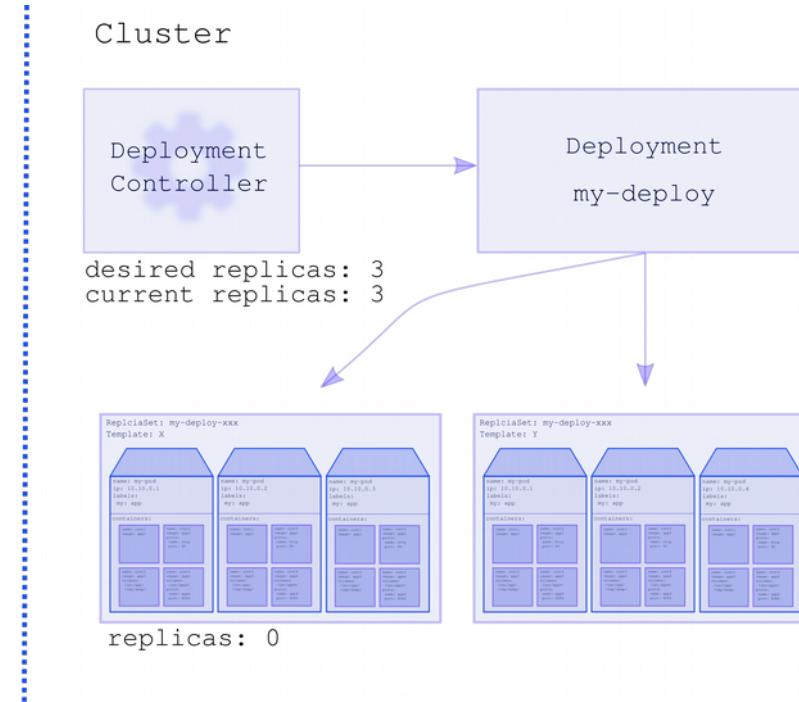
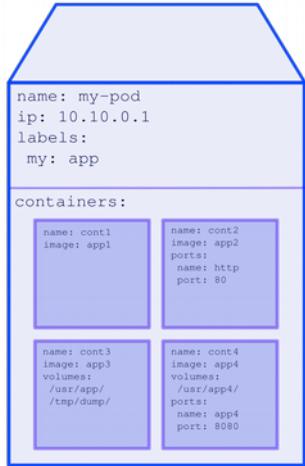


An update is triggered every time anything changes in the Pod template.

# Deployment



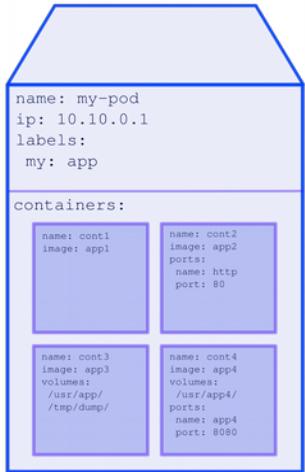
```
name: my-deploy  
replicas: 3  
strategy: RollingUpdate  
Template:
```



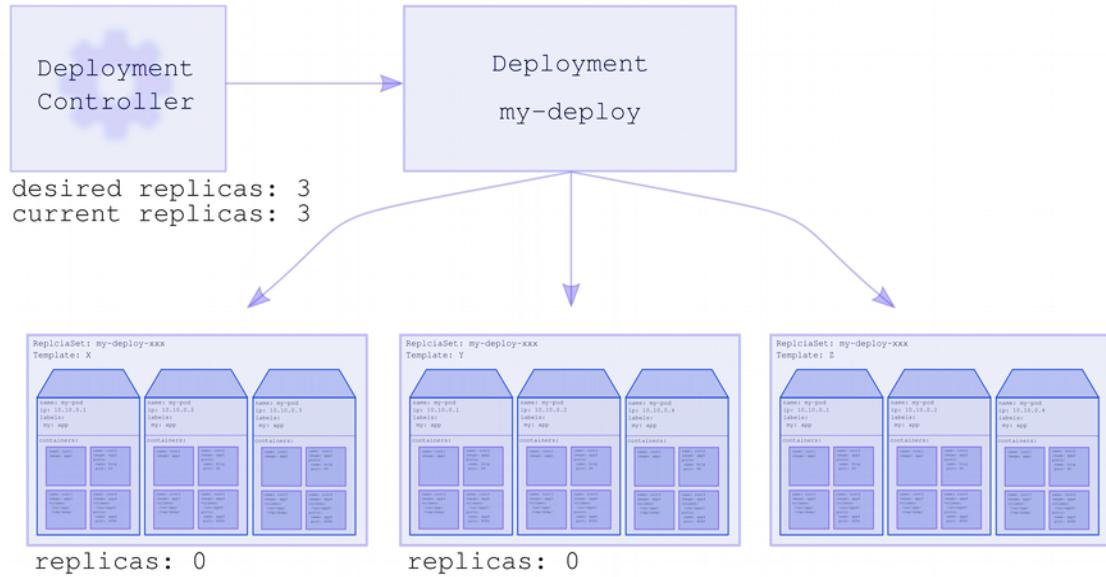
# Deployment



```
name: my-deploy  
replicas: 3  
strategy: RollingUpdate  
Template:
```



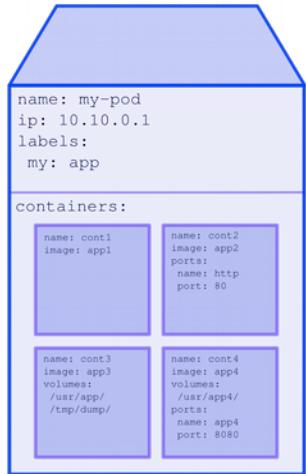
## Cluster



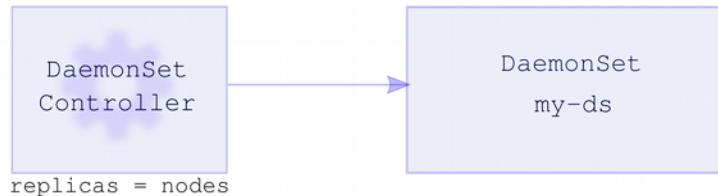
# DaemonSet



```
name: my-ds  
Template:
```



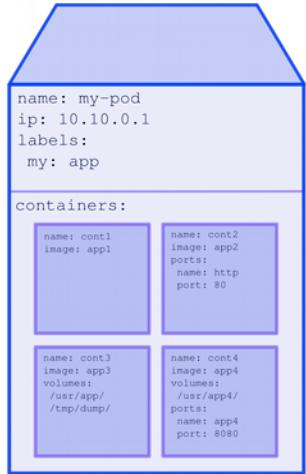
Cluster



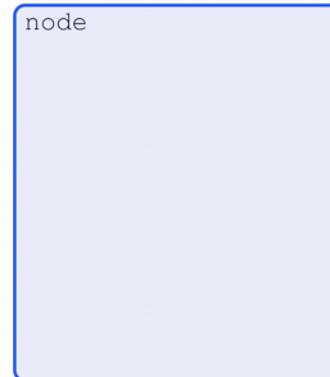
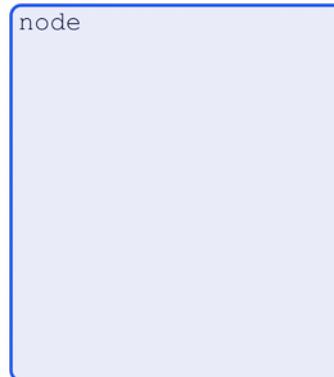
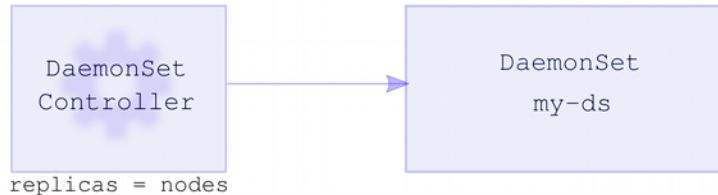
# DaemonSet



```
name: my-ds  
Template:
```



## Cluster



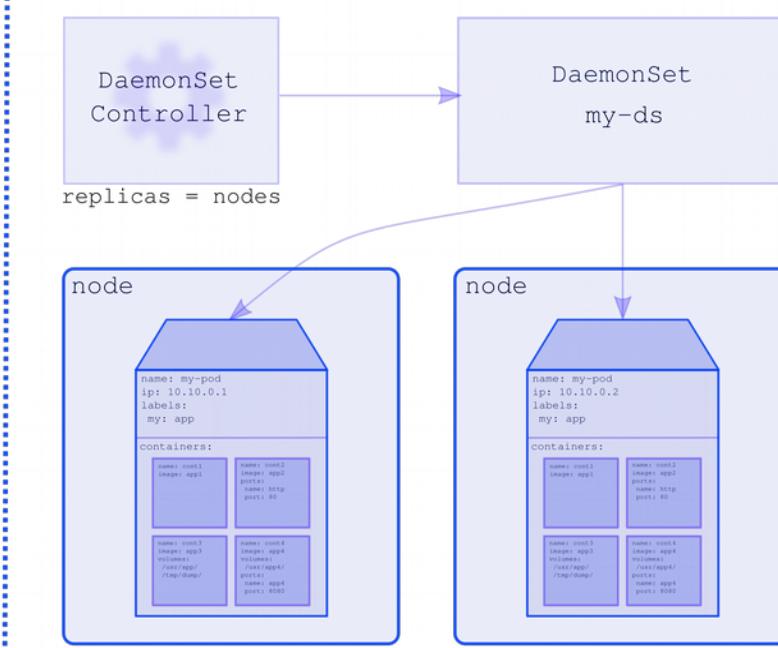
# DaemonSet



name: my-ds  
Template:



## Cluster



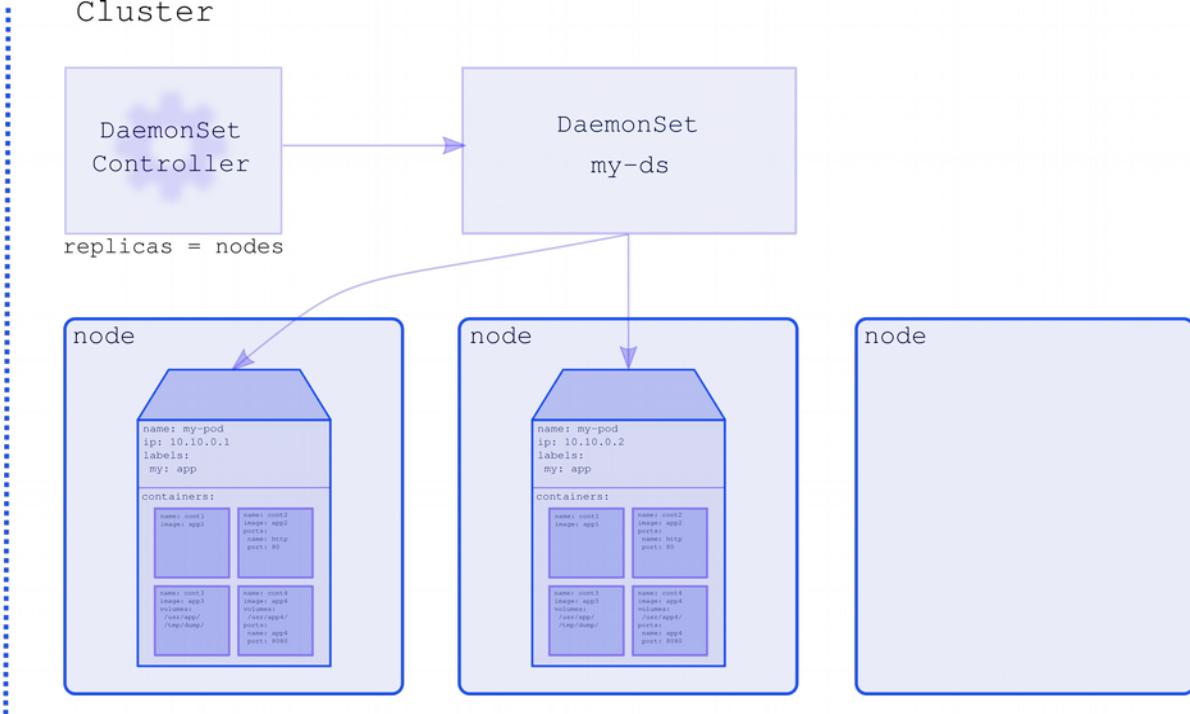
# DaemonSet



name: my-ds  
Template:



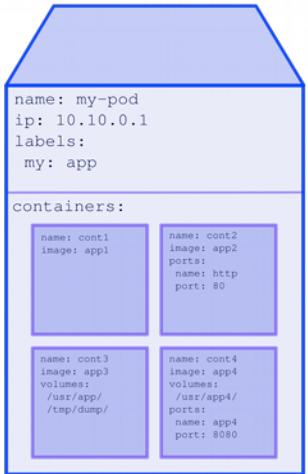
## Cluster



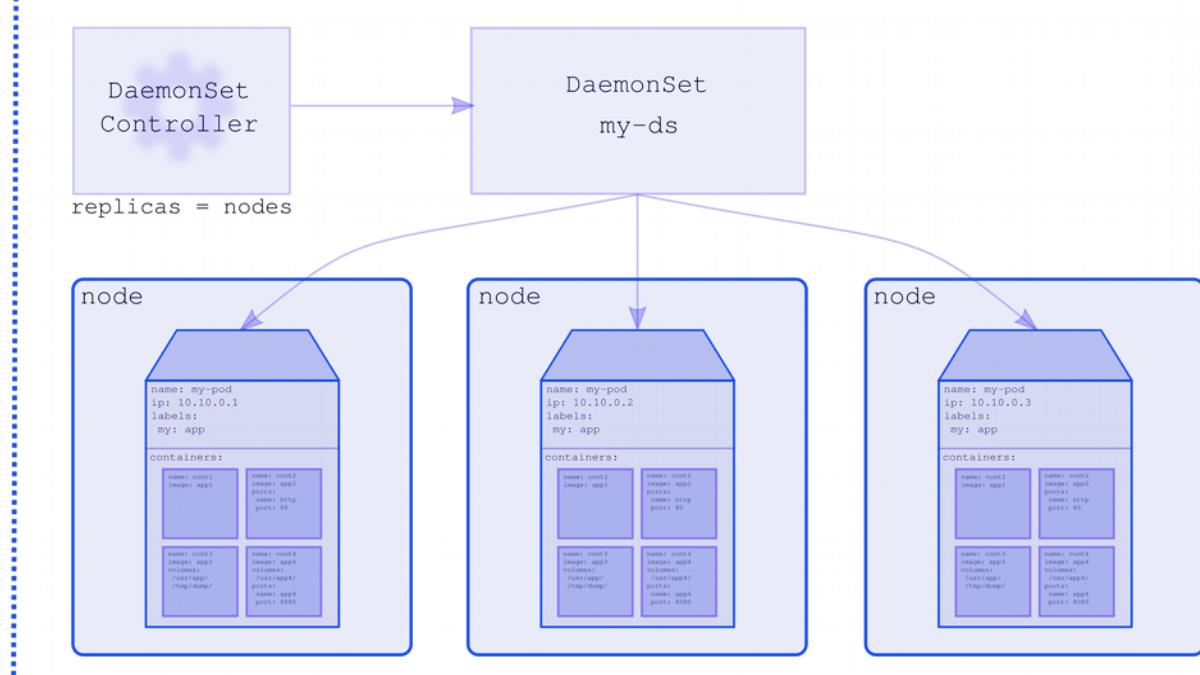
# DaemonSet



name: my-ds  
Template:



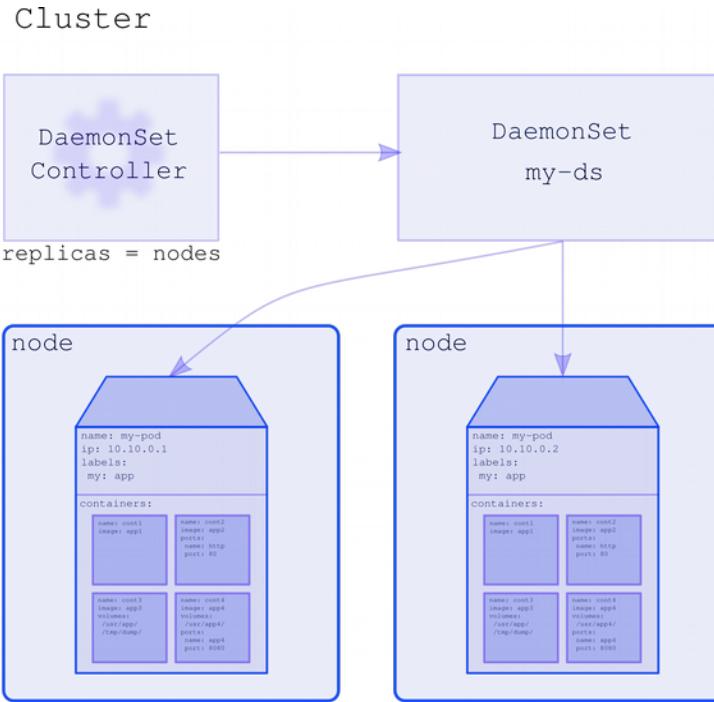
## Cluster



# DaemonSet



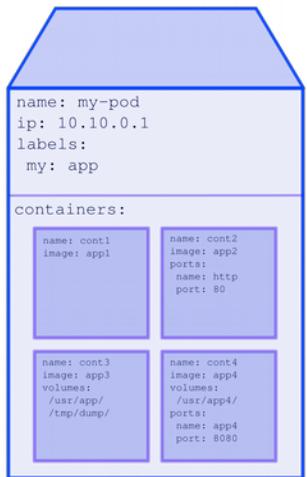
name: my-ds  
Template:



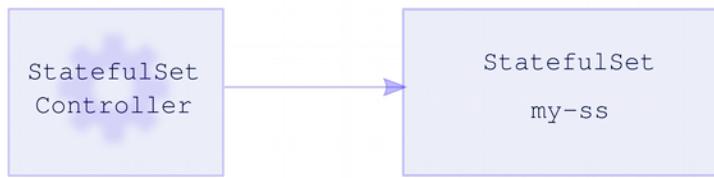
# StatefulSet



```
name: my-ss  
replicas: 3  
Template:
```



Cluster



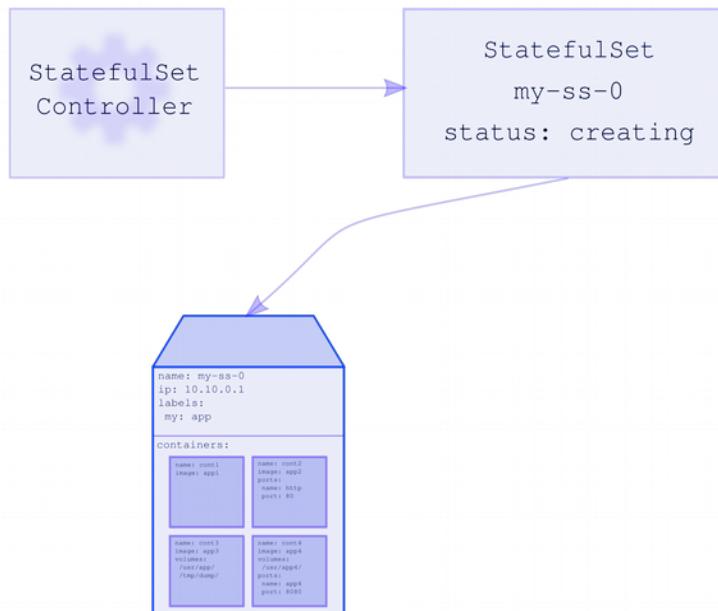
# StatefulSet



```
name: my-ss  
replicas: 3  
Template:
```



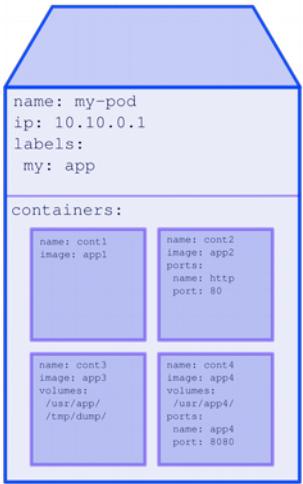
Cluster



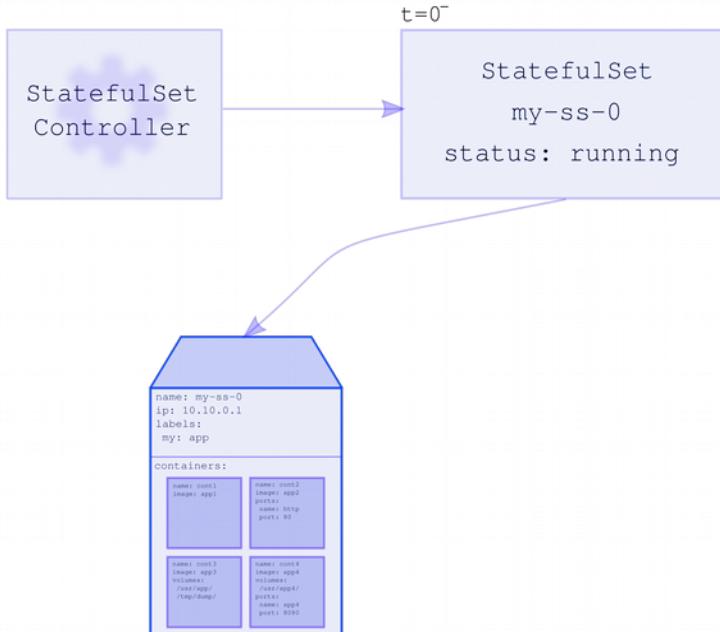
# StatefulSet



```
name: my-ss  
replicas: 3  
Template:
```



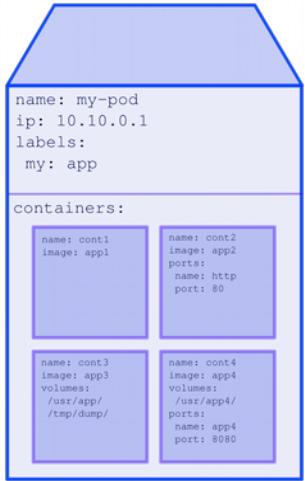
Cluster



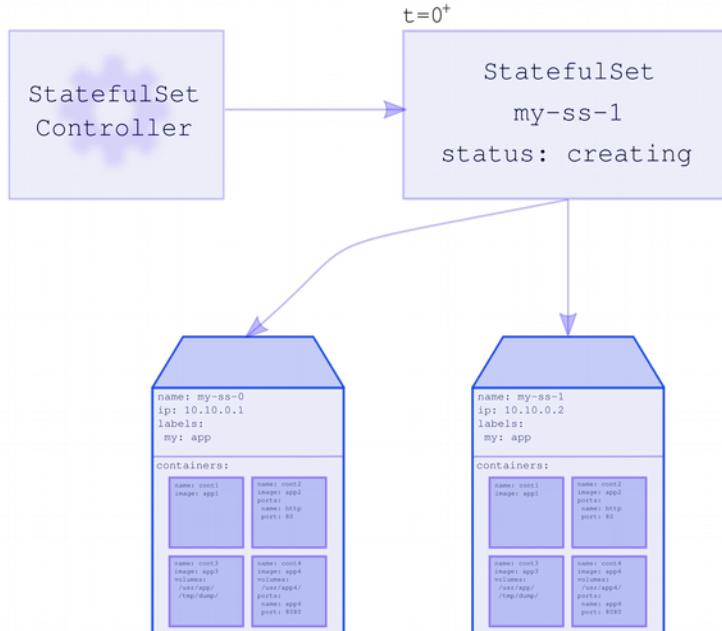
# StatefulSet



```
name: my-ss  
replicas: 3  
Template:
```



Cluster



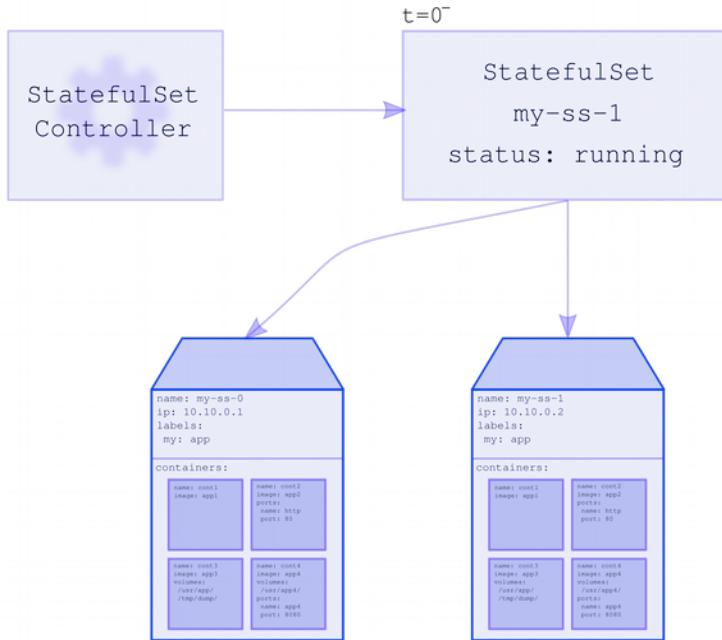
# StatefulSet



```
name: my-ss  
replicas: 3  
Template:
```



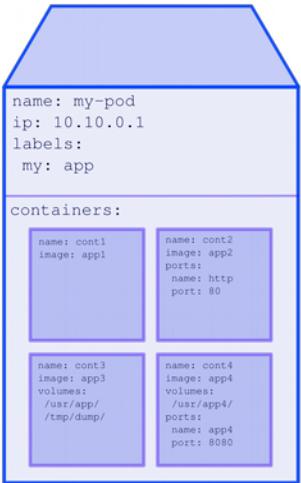
Cluster



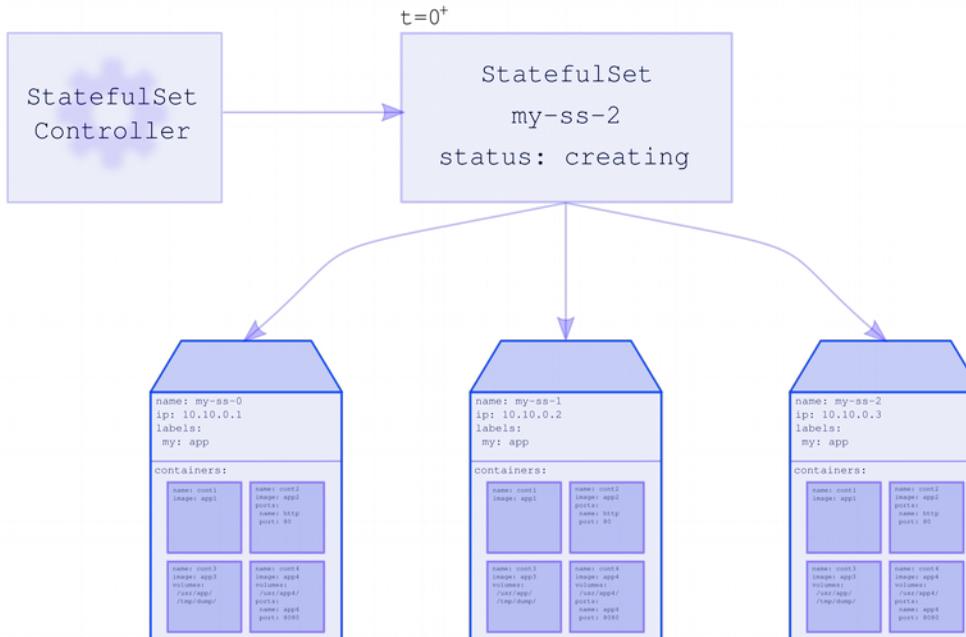
# StatefulSet



```
name: my-ss  
replicas: 3  
Template:
```



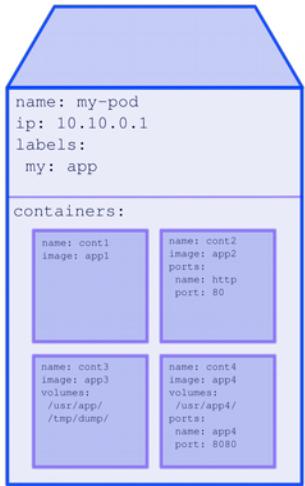
## Cluster



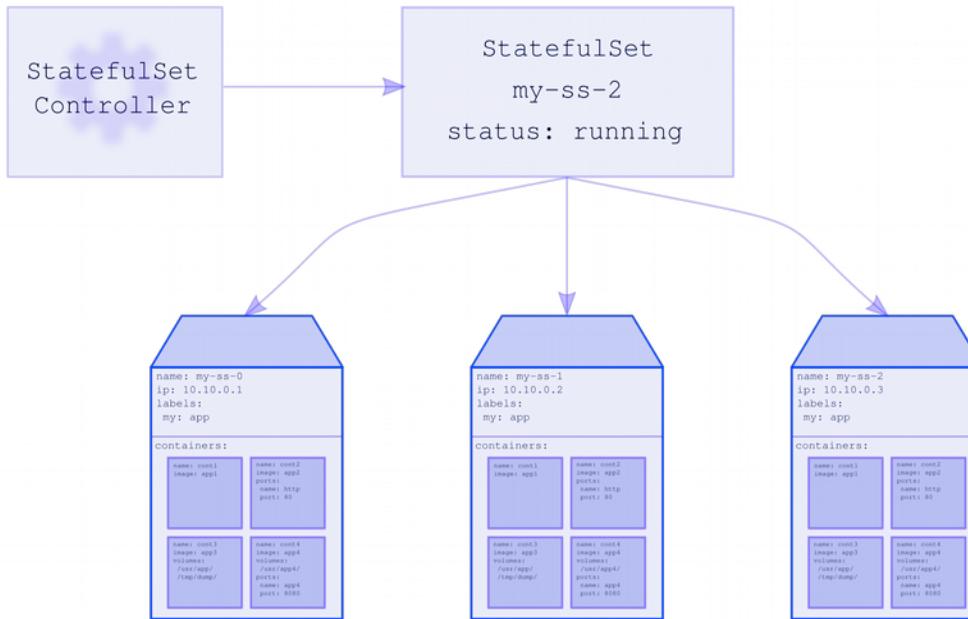
# StatefulSet



```
name: my-ss  
replicas: 3  
Template:
```



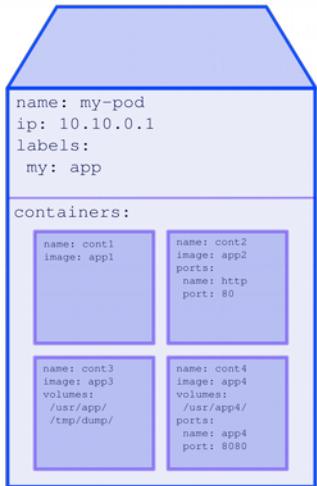
## Cluster



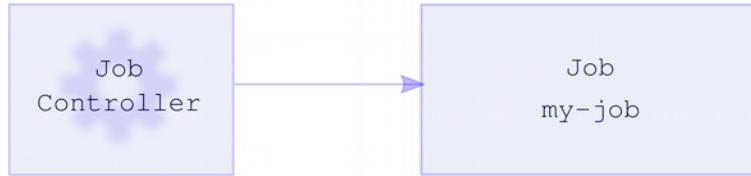
# Job/CronJob



```
name: my-job
backoffLimit: 5
activeDeadlineSeconds: 60
completions: 1
parallelism: 1
Template:
```



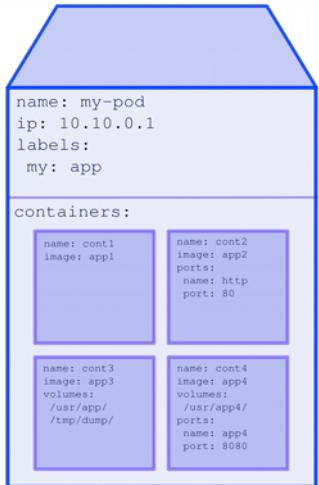
## Cluster



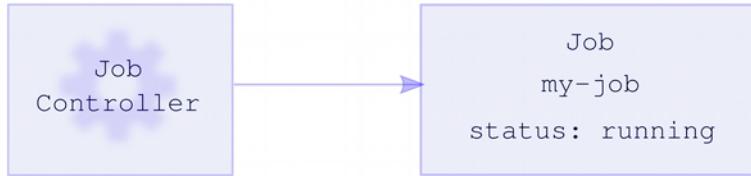
# Job/CronJob



```
name: my-job
backoffLimit: 5
activeDeadlineSeconds: 60
completions: 1
parallelism: 1
Template:
```



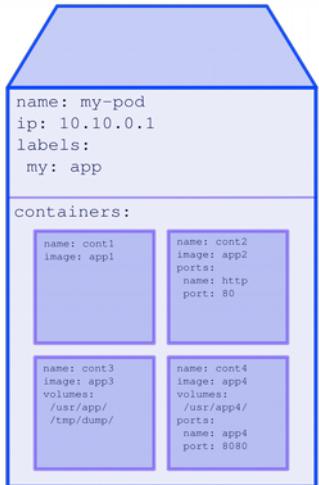
## Cluster



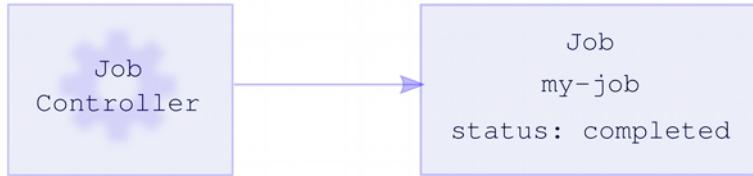
# Job/CronJob



```
name: my-job
backoffLimit: 5
activeDeadlineSeconds: 60
completions: 1
parallelism: 1
Template:
```



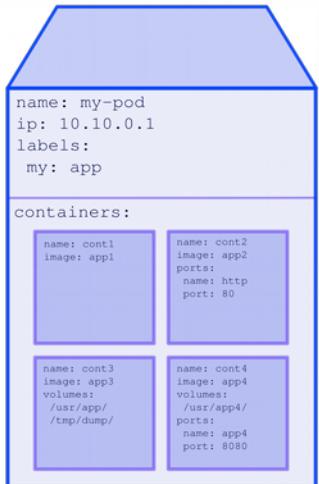
## Cluster



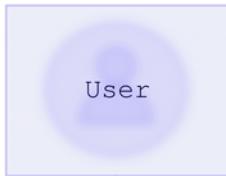
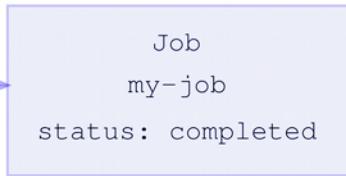
# Job/CronJob



```
name: my-job
backoffLimit: 5
activeDeadlineSeconds: 60
completions: 1
parallelism: 1
Template:
```



Cluster



logs

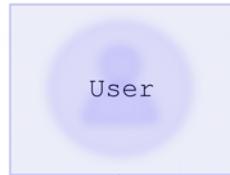
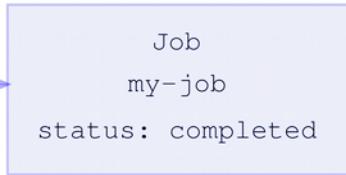
# Job - non-parallel



```
name: my-job
backoffLimit: 5
activeDeadlineSeconds: 60
completions: 1
parallelism: 1
Template:
```



Cluster

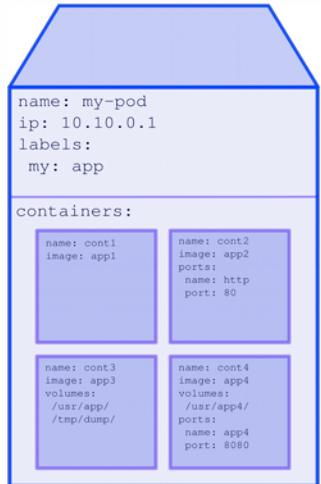


logs

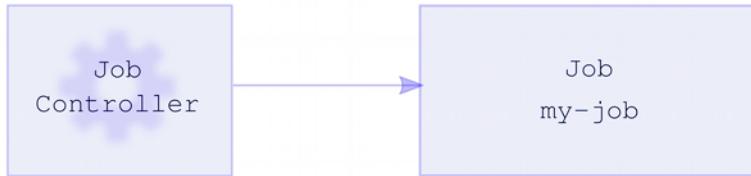
# Job - parallel, with fixed completion



```
name: my-job
backoffLimit: 5
activeDeadlineSeconds: 60
completions: X
parallelism: Y
Template:
```



Cluster



Job Controller will spin up one Job; with Y pods in parallel, when running in this mode

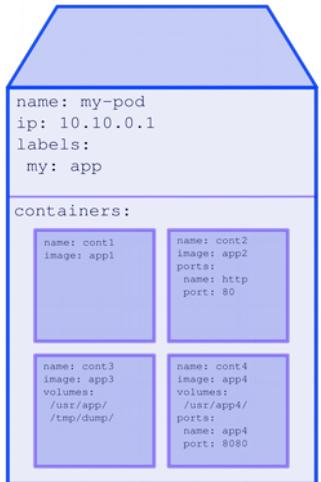


The Job will be considered finished, when X pods successfully complete.

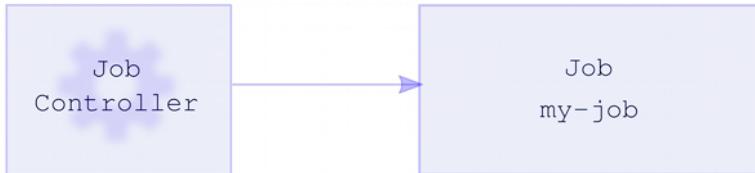
# Job - parallel, with work queue



```
name: my-job
backoffLimit: 5
activeDeadlineSeconds: 60
completions: null
parallelism: > 0
Template:
```



Cluster



To run in this mode,  
it requires running a  
queue service.

Each pod is independent  
to others.

Several outcomes are  
possible (depending on  
the app).

# Lab4

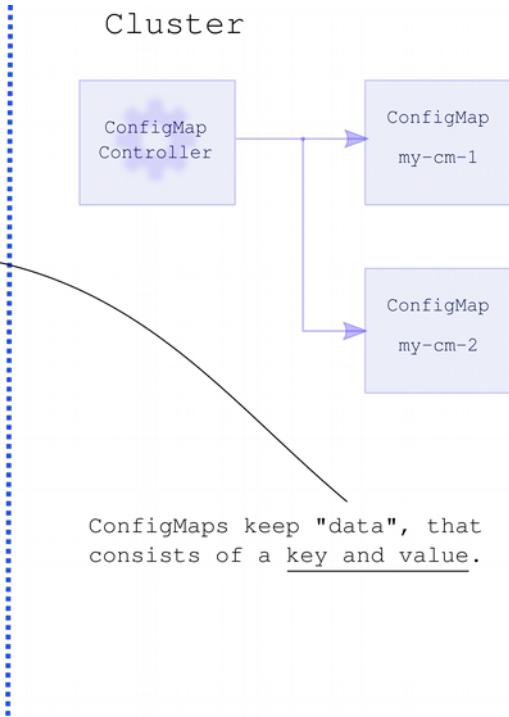


- 1.- Create a Deployment with deploy.yaml file.  
\$ kubectl create -f deploy.yaml
- 2.- Create a LoadBalancer type service with svc-lb.yaml.  
\$ kubectl create -f svc-lb.yaml  
Note: Observe the particularity of the service.
- 3.- Access each pod, go to /usr/share/nginx/html, and change the content of indeh.html to "pod-1" and "pod-2".  
\$ kubectl exec -it POD bash  
\$ echo "pod-X" > /usr/share/nginx/html/index.html
- 4.- SSH into the VM called lab4  
\$ gcloud compute ssh lab4 --zone ZONE
- 5.- Run a for loop with 20 requests to the IP address of the load balancer.  
\$ for i in {1..20}; \  
do curl IP\_ADDR; \  
done;
- 6.- Delete all the pods  
\$ kubectl delete pod --all  
Note: Observe K8s reaction
- 7.- To clean up, delete the deployment  
\$ kubectl delete deploy nginx
  
- 8.- Create a Job, with job.yaml  
\$ kubectl create -f job.yaml
- 9.- When the Job has completed, check the logs  
\$ kubectl logs POD  
Note: What do you think it is?

# ConfigMap



```
name: my-cm-1
data:
  CITY: BARCELONA
  ACKSTORM: TRUE
---
name: my-cm-2
data:
  file: my_data
```



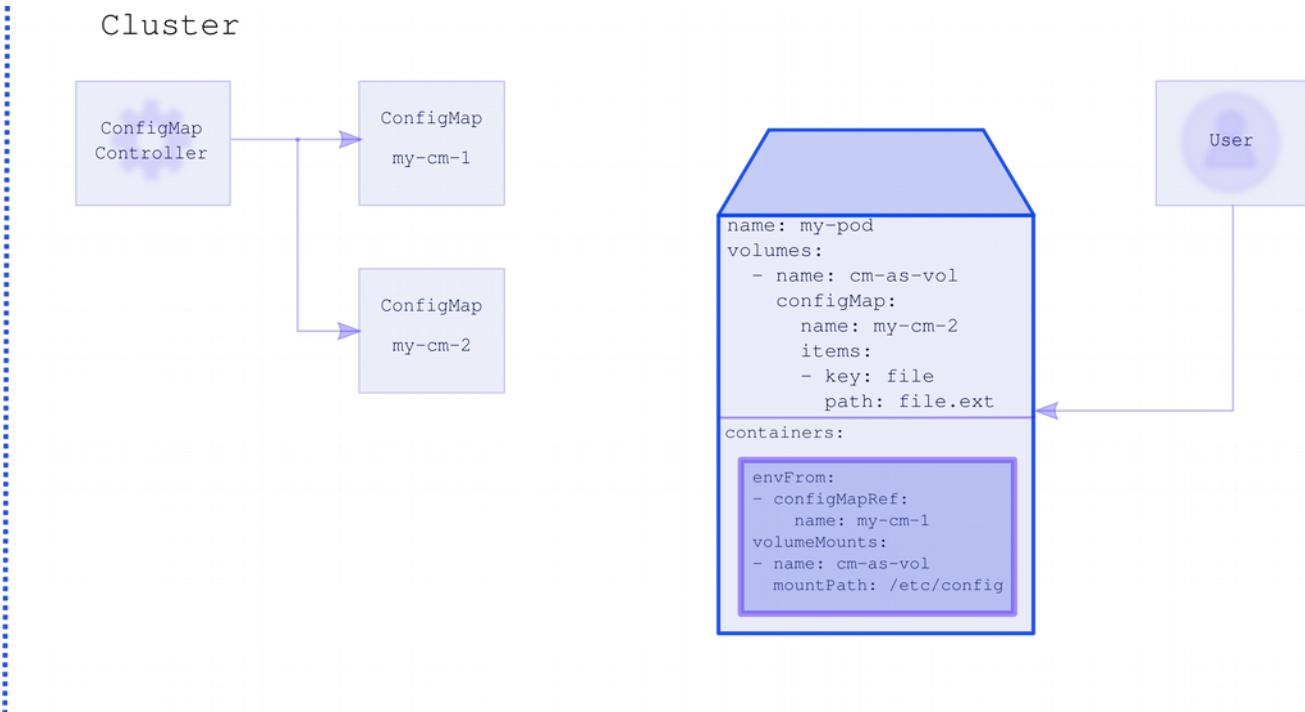
ConfigMaps allow to decouple config artifacts from image content to keep containerized applications portable.

ConfigMaps keep "data", that consists of a key and value.

# ConfigMap



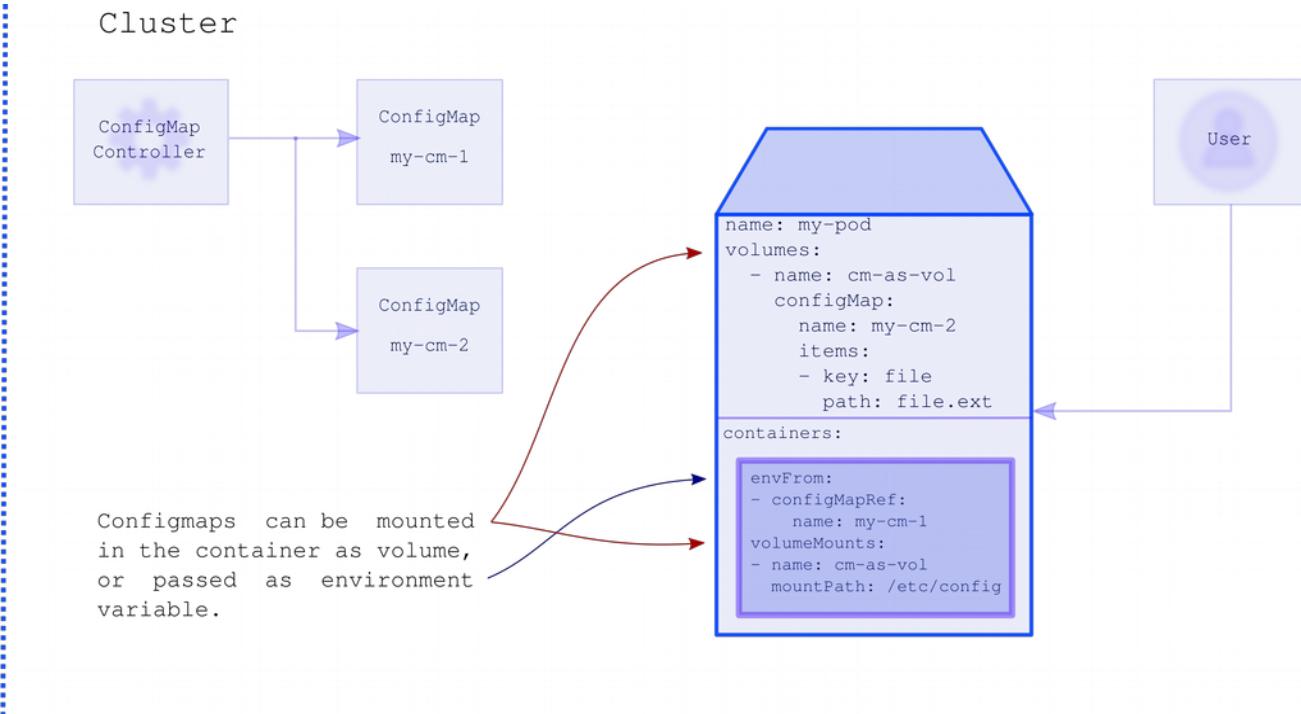
```
name: my-cm-1
data:
  CITY: BARCELONA
  ACKSTORM: TRUE
---
name: my-cm-2
data:
  file: my_data
```



# ConfigMap



```
name: my-cm-1
data:
  CITY: BARCELONA
  ACKSTORM: TRUE
---
name: my-cm-2
data:
  file: my_data
```

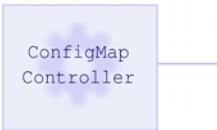


# ConfigMap

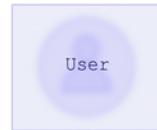


```
name: my-cm-1
data:
  CITY: BARCELONA
  ACKSTORM: TRUE
---
name: my-cm-2
data:
  file: my_data
```

Cluster



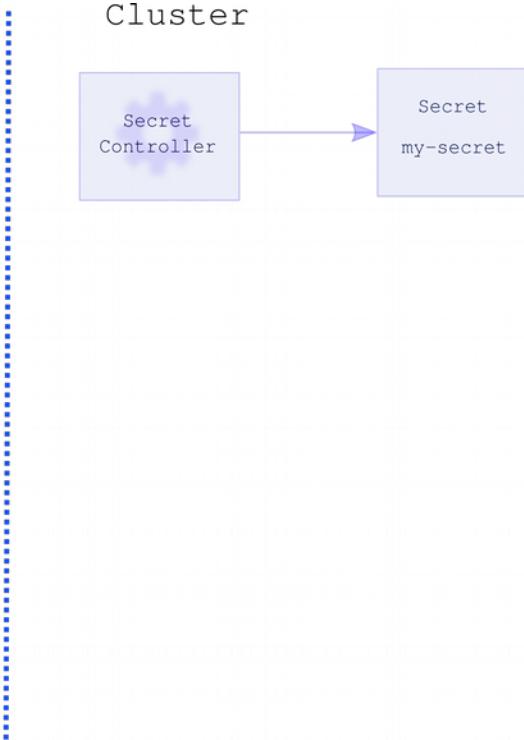
The ConfigMaps must exist when creating a pod that is going to use them.



# Secret



```
name: my-secret  
data:  
  username: user  
  password: pass
```



Secret objects allow to store and manage sensitive information, such as passwords, OAuth tokens, and ssh keys.

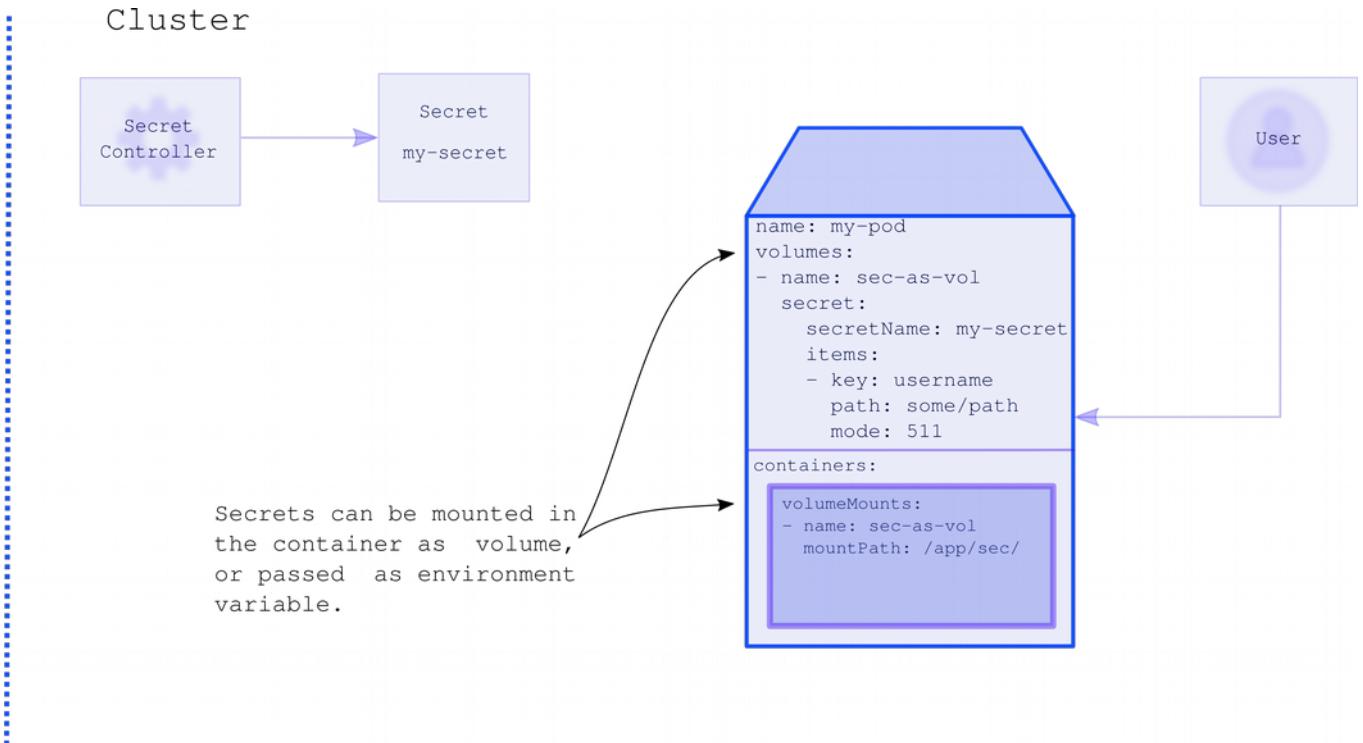
Similar to ConfigMaps, Secrets too allow decoupling some data from the application.

By decoupling sensitive data from the image or Pod specifications, we gain more control over how it is used.

# Secret



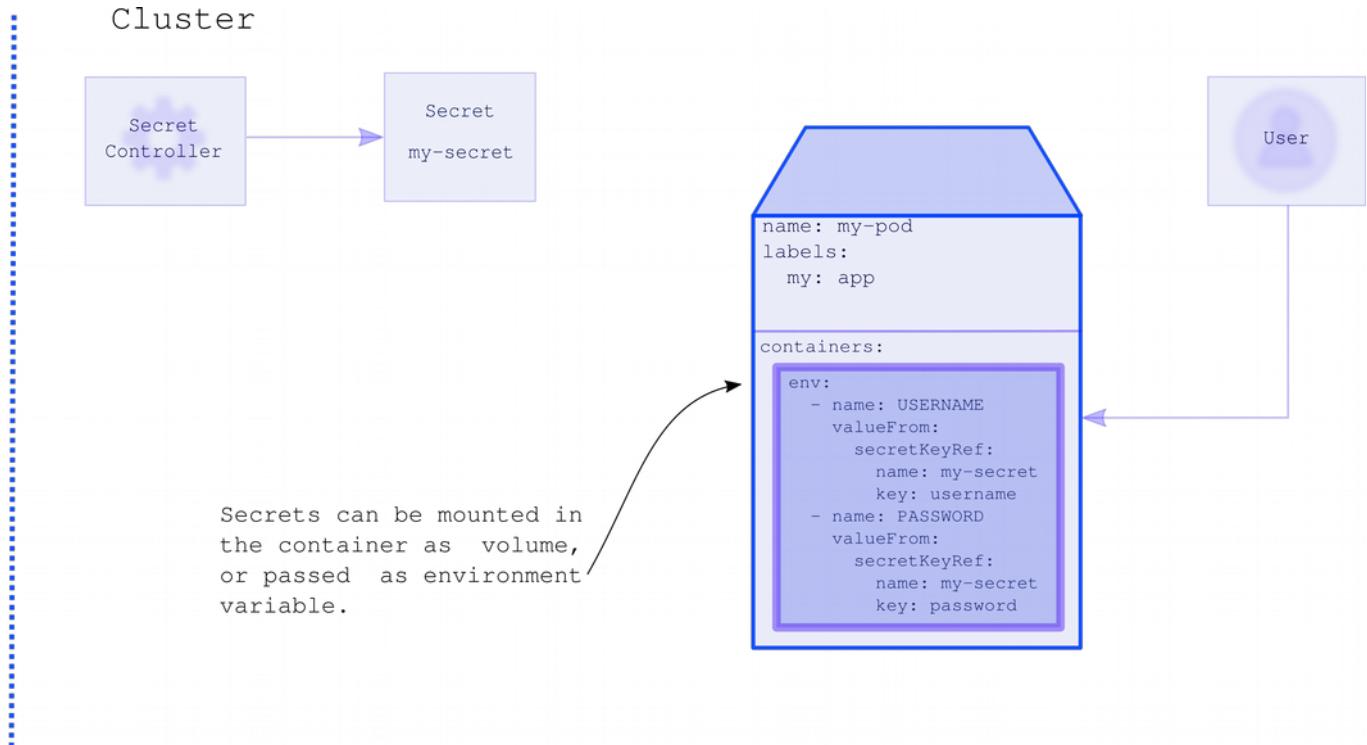
```
name: my-secret  
data:  
  username: user  
  password: pass
```



# Secret



```
name: my-secret  
data:  
  username: user  
  password: pass
```

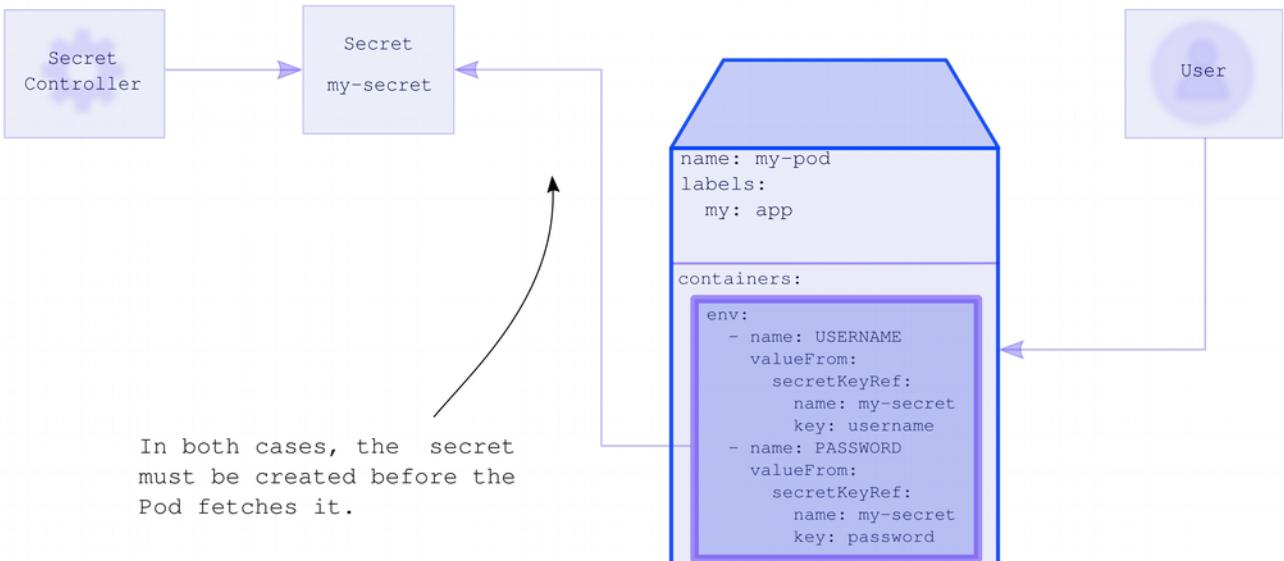


# Secret



```
name: my-secret  
data:  
  username: user  
  password: pass
```

## Cluster



# Lab4

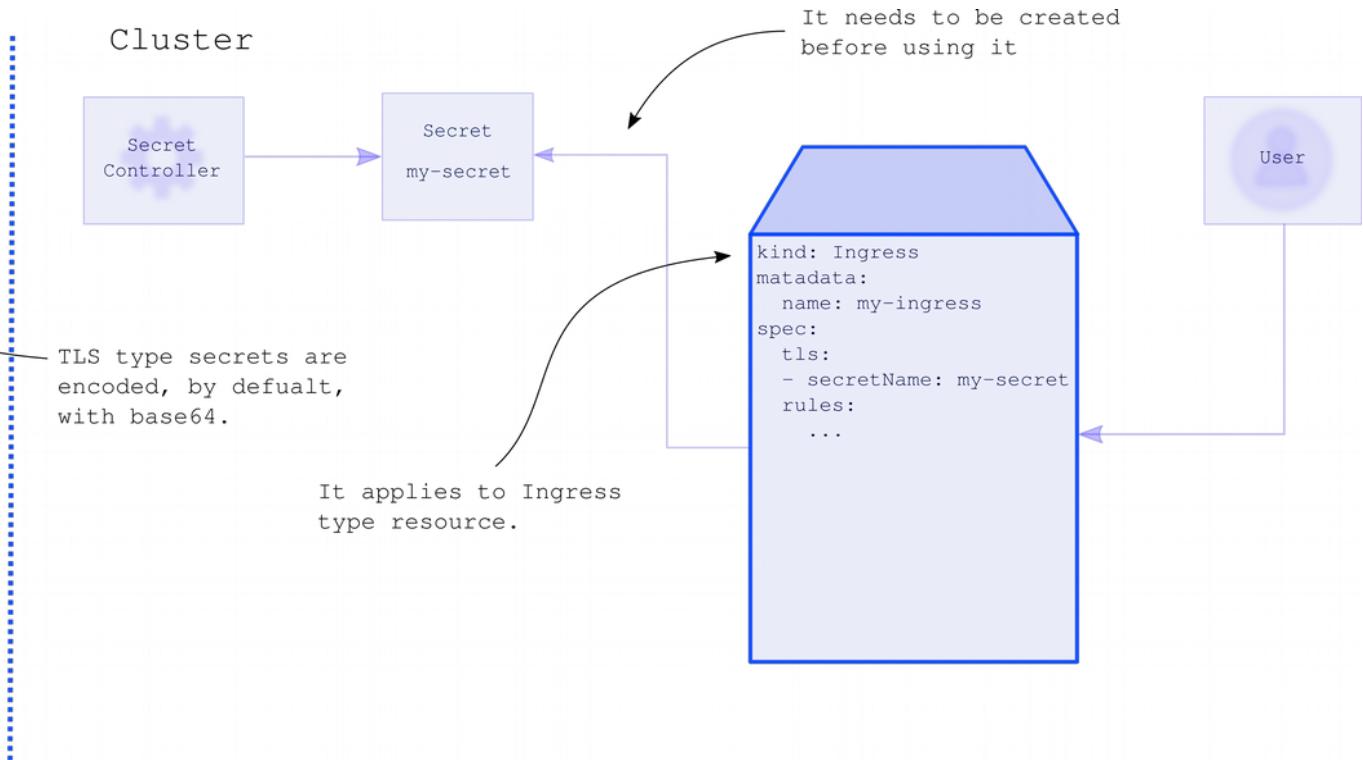


- 1.- Create a Deployment with deploy.yaml file.  
\$ kubectl create -f deploy.yaml
- 2.- Create a LoadBalancer type service with svc-lb.yaml.  
\$ kubectl create -f svc-lb.yaml  
Note: Observe the particularity of the service.
- 3.- Access each pod, go to /usr/share/nginx/html, and change the content of indeh.html to "pod-1" and "pod-2".  
\$ kubectl exec -it POD bash  
\$ echo "pod-X" > /usr/share/nginx/html/index.html
- 4.- SSH into the VM called lab4  
\$ gcloud compute ssh lab4 --zone ZONE
- 5.- Run a for loop with 20 requests to the IP address of the load balancer.  
\$ for i in {1..20}; \  
do curl IP\_ADDR; \  
done;
- 6.- Delete all the pods  
\$ kubectl delete pod --all  
Note: Observe K8s reaction
- 7.- To clean up, delete the deployment  
\$ kubectl delete deploy nginx
  
- 8.- Create a Job, with job.yaml  
\$ kubectl create -f job.yaml
- 9.- When the Job has completed, check the logs  
\$ kubectl logs POD  
Note: What do you think it is?

# Secret - TLS



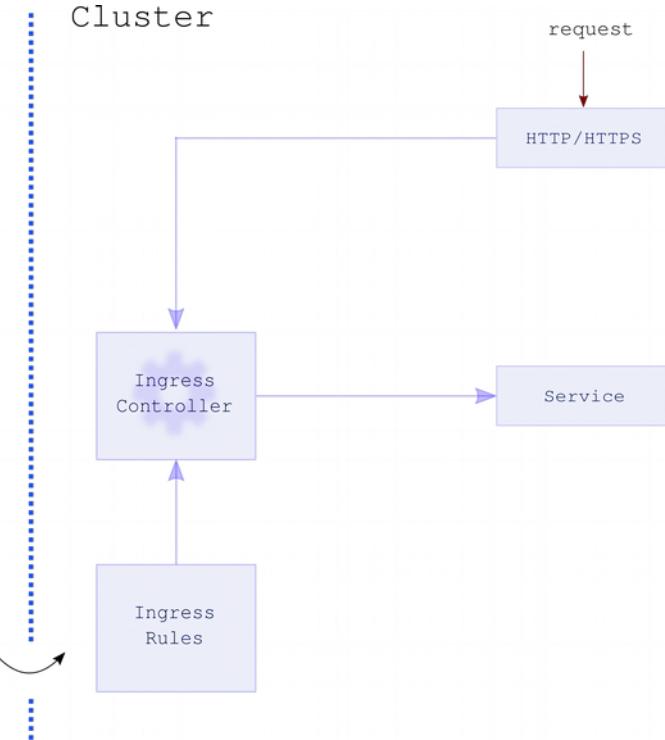
```
name: my-secret
type: tls
data:
  cert: cert | base64
  key: key | base64
```



# Ingress



```
name: my-ingress
spec:
  tls:
    - secretName: tls
  rules:
    - host: sitel.domain.com
      http:
        paths:
          - path: /path1/
            backend:
              serviceName: service1
              servicePort: 80
          - path: /path2/
            backend:
              serviceName: service2
              servicePort: 8080
          - path: /
            backend:
              serviceName: service3
              servicePort: 80
    - host: site2.domain.com
      http:
        paths:
          - backend:
              serviceName: service1
              servicePort: 80
```



Ingress object manages external access to the services in a cluster.

Traffic routing is controlled by rules defined on Ingress resource.

Ingress resource has no knowledge about the backends. It forwards a packet that matched an Ingress rule, to the right service.

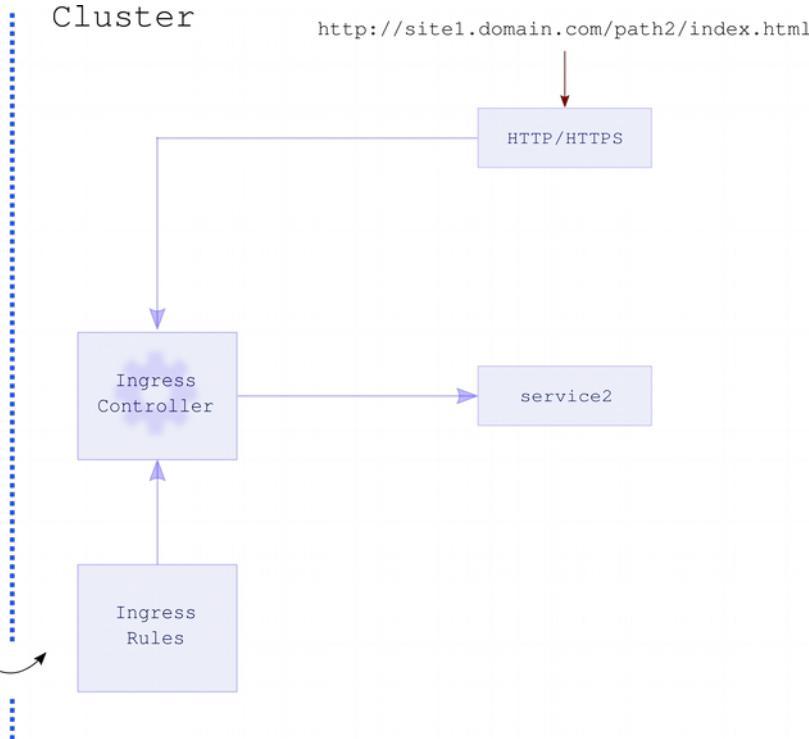
Ingress can provide load balancing, SSL termination and name-based virtual hosting.

Ingress Controller, unlike other controllers, is NOT part of the kube-controller-manager binary.

# Ingress



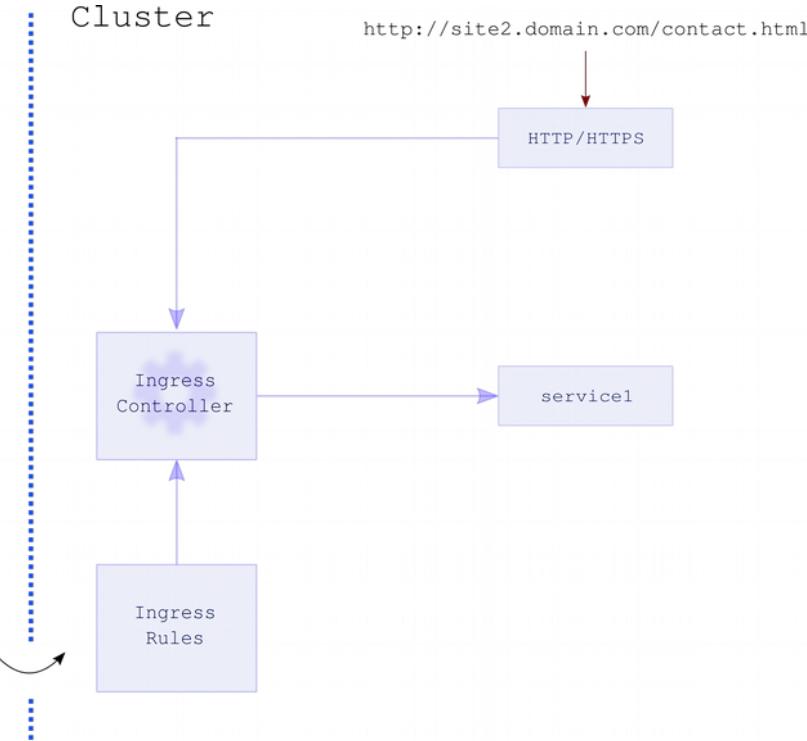
```
name: my-ingress
spec:
  tls:
    - secretName: tls
  rules:
    - host: site1.domain.com
      http:
        paths:
          - path: /path1/
            backend:
              serviceName: service1
              servicePort: 80
          - path: /path2/
            backend:
              serviceName: service2
              servicePort: 8080
          - path: /
            backend:
              serviceName: service3
              servicePort: 80
    - host: site2.domain.com
      http:
        paths:
          - backend:
              serviceName: service1
              servicePort: 80
```



# Ingress



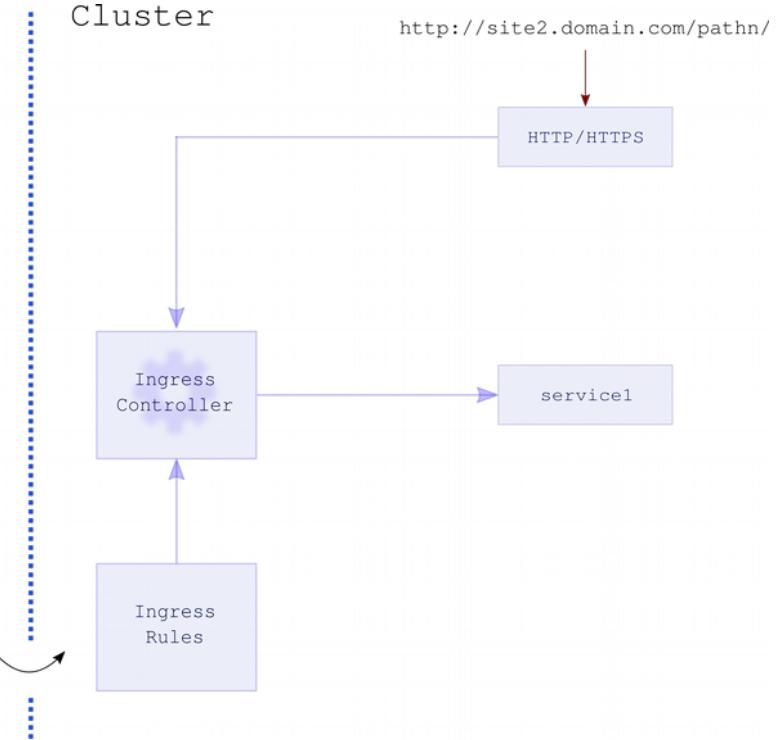
```
name: my-ingress
spec:
  tls:
    - secretName: tls
  rules:
    - host: site1.domain.com
      http:
        paths:
          - path: /path1/
            backend:
              serviceName: service1
              servicePort: 80
          - path: /path2/
            backend:
              serviceName: service2
              servicePort: 8080
          - path: /
            backend:
              serviceName: service3
              servicePort: 80
    - host: site2.domain.com
      http:
        paths:
          - backend:
              serviceName: service1
              servicePort: 80
```



# Ingress



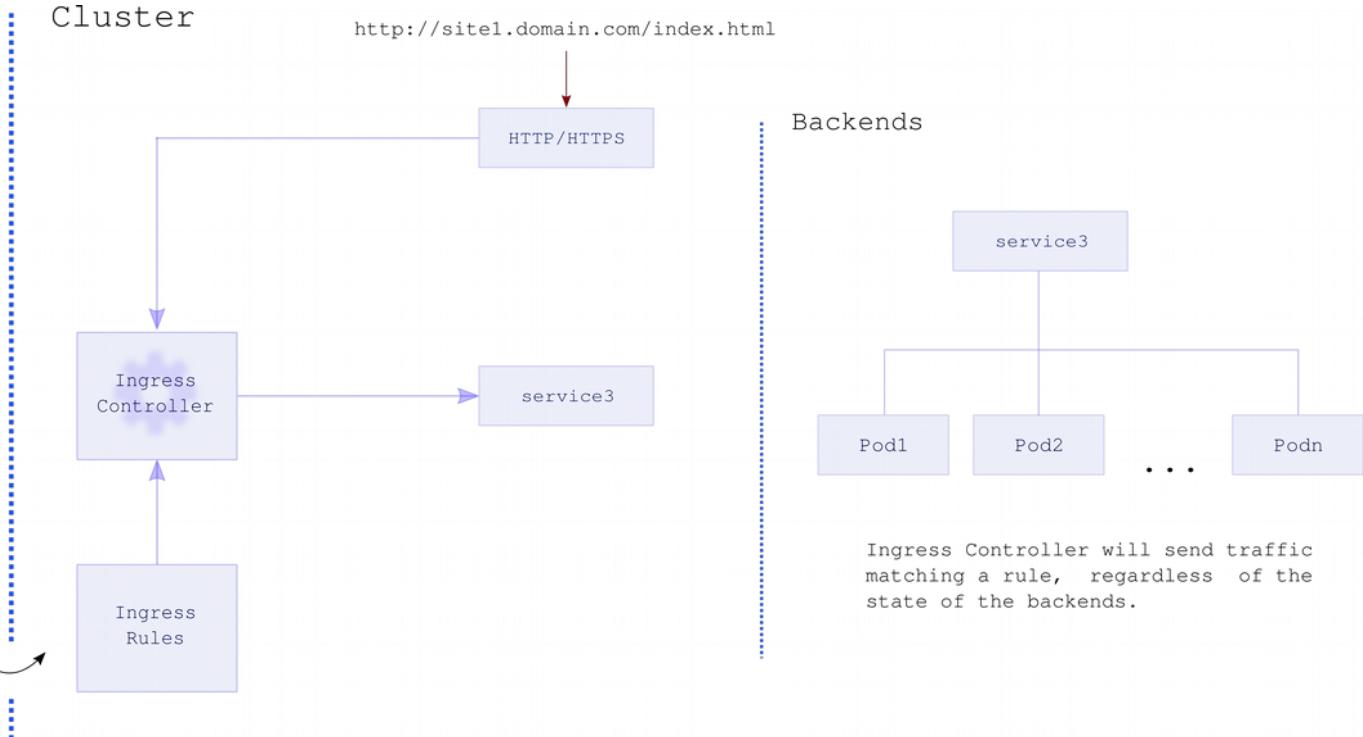
```
name: my-ingress
spec:
  tls:
    - secretName: tls
  rules:
    - host: site1.domain.com
      http:
        paths:
          - path: /path1/
            backend:
              serviceName: service1
              servicePort: 80
          - path: /path2/
            backend:
              serviceName: service2
              servicePort: 8080
          - path: /
            backend:
              serviceName: service3
              servicePort: 80
    - host: site2.domain.com
      http:
        paths:
          - backend:
              serviceName: service1
              servicePort: 80
```



# Ingress



```
name: my-ingress
spec:
  tls:
    - secretName: tls
  rules:
    - host: site1.domain.com
      http:
        paths:
          - path: /path1/
            backend:
              serviceName: service1
              servicePort: 80
          - path: /path2/
            backend:
              serviceName: service2
              servicePort: 8080
          - path: /
            backend:
              serviceName: service3
              servicePort: 80
    - host: site2.domain.com
      http:
        paths:
          - backend:
              serviceName: service1
              servicePort: 80
```



# Lab 5



Configure Redis login with a ConfigMap, and pass the password in a secret.

- 1.- Create the ConfigMap with cm.yaml file  
    \$ kubectl create -f cm.yaml
- 2.- Create the secret with secret.yaml file  
    \$ kubectl create -f secret.yaml
- 3.- Create the Redis pod, with pod.yaml file  
    \$ kubectl create -f pod.yaml
- 4.- Execute "ping" command on redis container.  
    \$ kubectl exec redis -- sh -c 'redis-cli ping'  
    Note: The command should fail with "NOAUTH Authentication required."
- 5.- Execute the same command, but now with the password; passed as an environment variable from the secret.  
    \$ kubectl exec redis -- sh -c 2>/dev/null 'redis-cli -a \$REDIS\_PASSWORD ping'

Create a deployment with certain pods targeted by an HTTPS Load Balancer, secured with TLS type secret.

- 1.- Create the landing page and nginx configuration ConfigMaps with cm-lp.yaml and cm-nginx.yaml
- 2.- Create 4 Nginx pods, and a NodePort type service targeting them.
- 3.- Create the certificates to be applied to the Load Balancer  
    \$ openssl req -newkey rsa:2048 -nodes -keyout tls.key -out tls.csr  
    \$ openssl x509 -in tls.csr -out tls.pem -req -signkey tls.key
- 4.- Create the secret from the certificates  
    \$ kubectl create secret tls certs --cert tls.pem --key tls.key
- 5.- Create the Ingress object, with ing.yaml file.  
    \$ kubectl create -f ing.yaml
- 6.- Once the Load balancer is created, and with healthy backends, curl it.  
    \$ curl https://LB\_IP\_ADDR/
- 7.- Modify the Ingress rules (from ing.yaml file), to accept requests for "training.com" host.

end

